

EXPLAINING NONLINEAR CLASSIFICATION DECISIONS WITH DEEP TAYLOR DECOMPOSITIONS AN ANALYSIS*

MARCEL POMMER

Abstract. During the last decade Deep Neural Networks (DNNs) as well as other sophisticated machine learning models gained substantially on relevance, due to so far unreached performance in a variety of topics like image recognition, classification or natural language processing, to name only a few. Despite their great performance those, mostly non linear models, lack of one important aspect, the explainability of the results. Montavon et al. introduce in their paper *Explaining NonLinear Classification Decisions with Deep Taylor Decompositions* a new technique, the deep taylor decomposition, to map the relevance of the output on the input features, i.e. quantify the influence of each input variable on the output. They demonstrate the results on two image recognition data sets, the MNIST and the ILSVRC, creating heatmaps to display the relevance of the single pixel. I recreate the procedure in Python and apply the deep taylor decomposition to a non image recognition data set, the titanic dataset.

Key words. Explainability, Deep Neural Networks, Image Recognition

MSC codes. 62H35, 93B15

1. Introduction. The raise of machine learning, combined with steadily growing computational power revolutionized many so far hard to grasp tasks like image recognition in order to push the development of self driving cars, help to diagnose diseases or automate classification problems. One can think of a variety of different models like random forests, boosting or deep neural networks and maybe even of more applications in our daily lives, beginning with the advertisement displayed in our browser, or the netflix recommendations going to automated voice recognition and translation to communicate with people all over the world. Those new techniques became quite famous during the last decade due to their overperformance in nearly ever field, however their complexity makes them mathematically hard to understand and explain leading to one of their major drawbacks, missing explainability. The paper *Explaining NonLinear Classification Decisions with Deep Taylor Decompositions* [XXX] by Montavon et al. tries to tackle this problem by extending the explainability of deep neural networks using taylor expansion, resulting in a mapping of non negative relevance from the output to each input feature. The authors apply their technique to a variety of exmaples from the MNIST [XXX] dataset as well as the ILSVRC [XXX] dataset. In figure 1 one can see a tractable example, in which a neural network detects a ‘0’ while distracted by a ‘5’. We denote the neurons with x_i and the respective contributions with R_i , resulting in a graphical representation, i.e. a heatmap, indicating which pixels contribut with which intensity to the decision of the neural network. Montavon et al. focus on image recognition in their paper, but highlight, that the procedure can be broadcasted to any input space and feature set. In my analysis of the application of the algorithm I will focus on the titanic [XXX] dataset which is very tractable and easy to understand and interpret.

The paper is organized as follows. The main results are in 2, the application on a simple neural network and algorithms in section 3, examples on the MNIST data set and experiments on the titanic dataset in section 4, and the conclusion follows in section 5.

*Submitted to professor Gutyniok on the 01.07.2022




FIG. 1. *Example: Detecting 0 while distracting other numbers are in place with a neural network*

2. Main results. In this section I will describe the general idea of the deep taylor decomposition, present the definitions and theorems as well as some handy examples. Since the authors focus on image recognition in their analysis I will follow this methodology, however in section ?? I will transfer all results to a simple regression task. In the context of image classification, we define the d -dimensional input $x \in \mathbb{R}^d$, where the image pixels (p) can be represented as $x = \{x_p\}$. The function $f(x) : \mathbb{R}^d \rightarrow \mathbb{R}^+$ quantifies either the probability of an object in the picture or the quantity of the object in question. The aim of the deep taylor decomposition is to assign a relevance score $R_p(x)$ to each pixel p in the input space. The relevance score quantifies the explanatory power of each pixel, i.e. the higher the relevance score the more important was the pixel for the classification. If the result is plotted in an image or to say heatmap as displayed in figure 1 the pixels which led to the classification decision are highlighted. In practice some conditions can help to further define and understand the relevance score. In the context of heatmaps, but also in other cases, the authors state three definitions.

DEFINITION 2.1 (conservative). *A heatmapping $R(x)$ is conservative if the sum of assigned relevances in the pixel space corresponds to the total relevance detected by the model, that is*

$$\forall x : f(x) = \sum_p R_p(x)$$

In other words, the sum of the relevance of all pixels should align with the output, so the probability or quantity of the object in question. Coming back to figure 1, if the output of the neural network defines a probability of 90 % the sum of the relevance of all pixels should be 90 % or 0.9 as well. Definition 2.1 ensures that all relevance detected by the model can be explained by the input variables.

DEFINITION 2.2 (positiv). *A heatmapping $R(x)$ is positiv if all values forming the heatmap are greater or equal to zero, that is:*

$$\forall x, p : R_p(x) \geq 0$$

This property ensures, that relevance cannot be negative in a sense that two pixels cancel each other out. In other words, an input feature either has a positive impact on the decision made, or no impact at all, but there is no contradictory evidence.

Since definition 2.1 and definition ?? are of essence for the evaluation of models we further define:

DEFINITION 2.3 (consistent). *A heatmapping $R(x)$ is consistent if it is conservative and positive.*

We will use definition 2.3 to evaluate heatmaps, but it has to be mentioned that many relevance rules might confirm with the definition of *consistency*, however its not

a measure for the quality, which can be seen in the following example of a uniform relevance over all pixels:

$$\forall p : R_p(x) = \frac{1}{d}f(x),$$

where d denotes the number of pixels. Although, the heatmap will comply with definition 2.3 it will result in an all black image giving no further information on the relation between input and output.

3. Algorithms. The deep taylor decomposition is based on the first order taylor expansion at a root point \tilde{x} , such that $f(\tilde{x}) = 0$:

$$(3.1) \quad f(x) = f(\tilde{x}) + \left(\frac{\partial f}{\partial x} \Big|_{x=\tilde{x}} \right)^T \cdot (x - \tilde{x}) + \epsilon = 0 + \sum_p \frac{\partial f}{\partial x_p} \Big|_{x=\tilde{x}} \cdot (x_p - \tilde{x}_p) + \epsilon,$$

where the sum over all pixels derivative is defined as the redistributed relevance:

$$R(x) = \frac{\partial f}{\partial x} \Big|_{x=\tilde{x}} \odot (x - \tilde{x}),$$

and \odot is defined as element wise multiplication. The finding of the root point is a great challenge an far from obvious. In figure 2 we can see a picture of the monopterus in Munich, where the root point is simply a variant of the picture where the building is blurred. Since an image possibly can have more than one root point the choice is




FIG. 2. Example: Root point in an image.

crucial and a good root point deviates from the original point x as few as possible, i.e. minimizing the objective:

$$\min_{\xi} \|\xi - x\|^2 \text{ subject to } f(\xi) \text{ and } \xi \in \mathbb{X},$$

where \mathbb{X} is the input domain.

Next we focus on the deep taylor decomposition, where we consider the mapping of neurons in one layer to each neuron in the next layer, assuming a relation explainable by some relevance function $R_j(\{x_i\})$. If this assumption holds and we further can identify a root point $\{\tilde{x}_i\}$ such that $R_j(\{\tilde{x}_i\}) = 0$ we can write equation 3.1:

$$\sum_j R_j = \left(\frac{\partial(\sum_j R_j)}{\partial\{x_i\}} \Big|_{\{\tilde{x}_i\}} \right)^T \cdot (\{x_i\} - \{\tilde{x}_i\}) + \epsilon = \sum_i \sum_j \frac{\partial R_j}{\partial x_i} \Big|_{\{\tilde{x}_i\}} \cdot (x_i - \tilde{x}_i) + \epsilon$$

If definition 2.3 holds, the relevances is guaranteed to be conserved in each layer, i.e. $R_f = \dots = \sum_j R_j = \dots = \sum_p R_p$ and $R_f, \dots, \{R_j\}, \dots, \{R_p\} \geq 0$.

Further I will present two different approaches, the ω^2 -rule and the z -rule.

As a starting point we consider a simple detection pooling neural network with a reluctified linear activation function, i.e.

$$x_j = \max(0, \sum_i x_i w_{i,j} + b_j), \quad x_k = \sum_j x_j,$$

where $\{x_i\}$ is a d-dimensional input and $\theta = \{w_{i,j}, b_j\}$ are weights and bias. To guarantee the exististence of a root point in the origin $b_j \leq 0$. The relevance of the top layer is due to the pooling $R_k = \sum_j x_j$ and can be redistributed to the next layer according to the taylor decomposition

$$R_j = \left. \frac{\partial R_k}{\partial x_j} \right|_{\{\tilde{x}_j\}} \cdot (x_j - \tilde{x}_j)$$

If either $\sum_j \tilde{x}_j = 0$ and $\forall j : \tilde{x}_j \geq 0$ we need to chose $\{\tilde{x}_j\}$ resulting in $R_j = x_j$, since $\frac{\partial R_k}{\partial x_j} = 1$. Redistributing the relevance to the next layer using the taylor decomposition leads to

$$(3.2) \quad R_j = \sum_i \left. \frac{\partial R_j}{\partial x_i} \right|_{\{\tilde{x}_i\}^{(j)}} \cdot (x_i - \tilde{x}_i^{(j)})$$

which is the starting point for the further analysis.

3.1. Unconstrained Input Space and ω^2 -Rule. If we consider an unconstrained input space one can chose the nearest root point in the Euclidean sense. Considering the reluctified linear activation the intersection of the equation $\sum_i \tilde{x}_i^{(j)} w_{i,j} + b_j$ and the line of maximum descent defined by the derivative $\{\tilde{x}_i\}^{(j)} = \{x_i\} + t w_j$, where w_i ist the weight vector and $t \in \mathbb{R}$. The root point is then given by $\{\tilde{x}_i\}^{(j)} = \{x_i - \frac{w_{ij}}{\sum_i w_{ij}^2} (\sum_i x_i w_{ij} + b_j)\}$ which leads by pluggin in 3.2

$$R_i = \sum_j \frac{w_{ij}^2}{\sum_i w_{ij}^2} R_j$$

PROPOSITION 3.1 (w-Rule consistency). *$\forall g \in G$, the deep Taylor decomposition with the w^2 -rule is consistent in the sense of Definition 3.*

3.2. Constrained Input Space and the z-Rules. Since in many cases the input domain is restricted the authors present a rule for bounded input spaces as well. Montavon et. al consider in their paper two possible spaces, first $\mathbb{X} = \mathbb{R}_x^d$ and $\mathbb{B} = \{\{x_i\} : \forall_{i=1}^d l_i \leq x_i \leq h_i\}$, where l_i and h_i are the respective lower and higher bounds for each input feature. I will only cover the first case since logic and results are quiet similar and the domain corresponds to the rectified linear activation. Hence we already know of the existence of one root at the origin we search on the segment $(\{x_i 1_{w_{ij} < 0}\}, \{x_i\})$ and thus the direction of the segment is given by $v_i^{(j)} = x_i - x_i 1_{w_{ij} < 0} = x_i 1_{w_{ij} \geq 0}$. If we follow the same logic is in section but change the line of maximum descent to $\{\tilde{x}_i\}^{(j)} = \{x_i\} + t x_i 1_{w_{ij} \geq 0}$ we get the following relevance propagation rule

$$R_i = \sum_j \frac{x_i 1_{w_{ij} \geq 0}}{\sum_i x_i 1_{w_{ij} \geq 0}} R_j$$

PROPOSITION 3.2 (z-Rule consistency). *$\forall g \in G$, the deep Taylor decomposition with the z-rule is consistent in the sense of Definition 3.*

3.3. Deep Neural Networks. Since many neural networks use complex deep architectures the authors further show a tractable way for the mapping of relevance from higher to the lower layers and introduce the concept of relevance models. The Min-Max and the Training-Free relevance model are introduced in the paper, however I will only cover the first. It is defined as

$$y_j = \max(0, \sum_i x_i v_{ij} + a_j), \quad \hat{R}_k = \sum_j y_j,$$

where $a_j = \min(0, \sum_l R_l v_{lj} + d_j)$ is a negative bias where the sum runs over the detection neurons from the upper layer and R_l are the corresponding relevances. After estimation of the parameters $\{v_{ij}, v_{lj}, d_j\}$ by minimizing

$$(3.3) \quad \min \langle (\hat{R}_k - R_k)^2 \rangle$$

, where R_k and \hat{R}_k are the true and predicted relevances, we end up with the same problem as in section 2. Due to the similar structure we can apply the same computations and compute $R_j = y_j$ and $R_i = \sum_j \frac{q_{ij}}{\sum_i q_{ij}} R_j$, where $q_{ij} = \{v_{ij}^2, x_i v_{ij}\}$ for the w^2 -rule and z-rule, respectively. In contrast to the problem before the resulting heatmap is only approximately conservative, due to the possible errors during the minimization of 3.3.

4. Experiment. As mentioned in the beginning the authors focus on image recognition and present examples from the MNIST and the ILSVRC data sets. I will only present one example from the MNIST data set and conclude with a self made example from the titanic data set.

4.1. MNIST Example.

4.2. Titanic Example.

5. Conclusion. During the last years the importance of explainability of complex machine learning models became more and more present, due to either simple questions of understanding or even legal issues in for example pricing of insurances. Montavon et al. present in their paper *Explaining NonLinear Classification Decisions with Deep Taylor Decompositions* practical algorithms which can help to tackle the problem of interpretability for a wide range of deep learning models. the authors furthermore substantiate their theory with examples on two very famous image recognition data sets, namely the MNIST and the ILSVRC data sets. I myself could also reproduce very reliable results on the titanic data set which is in contrast to the other two a simple supervised binary classification problem.