Explaining NonLinear Classification Decisions with Deep Taylor Decomposition

Marcel Pommer

LMU München

12. Juni 2022

Explainability

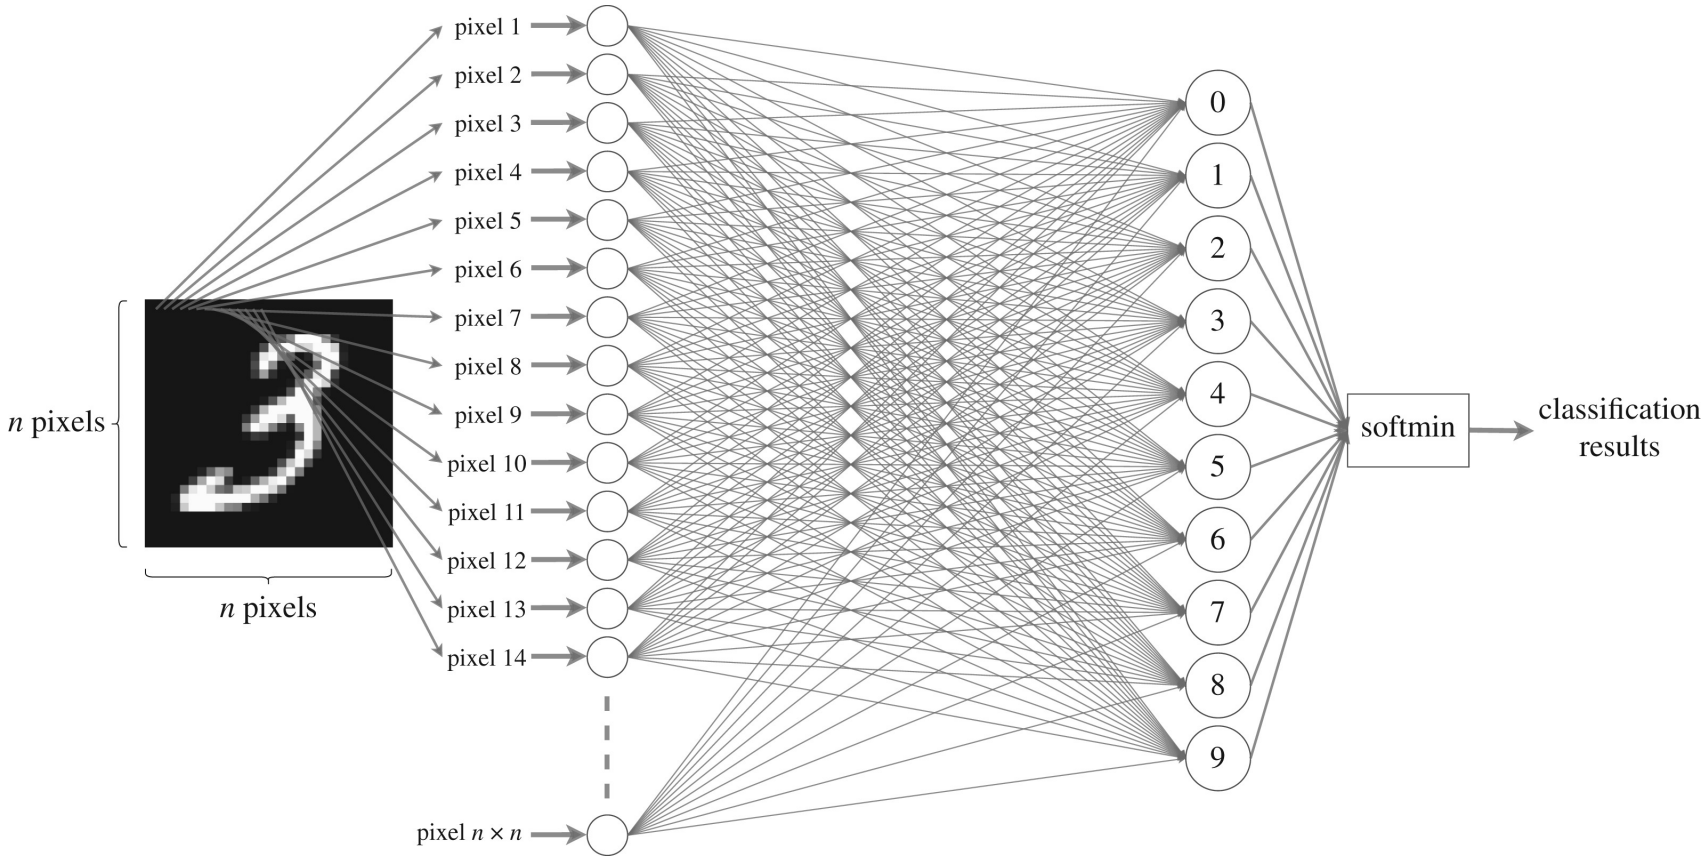

Deep neural networks performe great on a variety of problems
but how can we justify decisions made by complex deep architectures?[4]

Table of Contents

Table of Contents

Introduction

Deep neural networks revolutionized amongst others the field of

- Image recognition
- Natural language processing
- Human action recognition
- Physics
- Finance
- ...

With one major drawback $\rightarrow$ **lack of transparency**

Interpretable Classifier

Explanation of non-linear classification in terms of the inputs
$\rightarrow$ A classifier should not only provide a result but also a reasoning

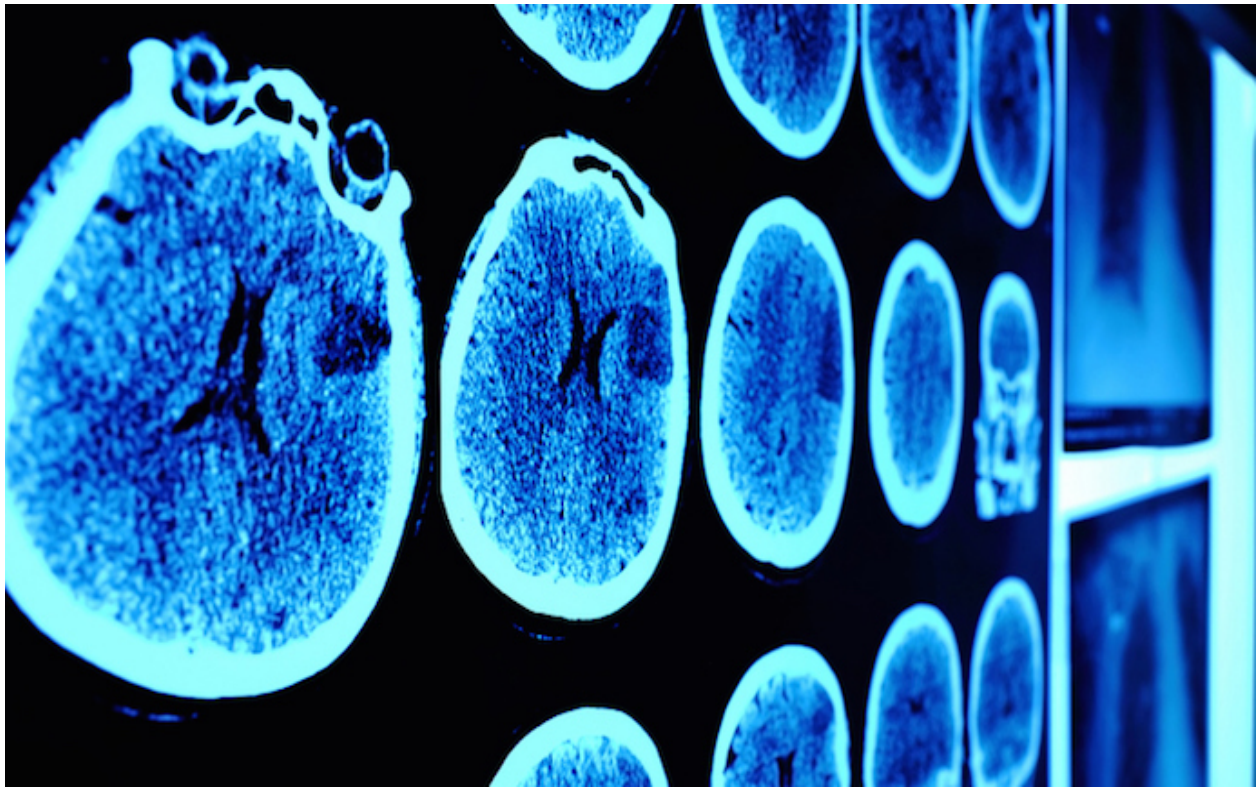

We do not only need to know if the patient has cancer but also where exactly it is located

General Idea

To accomplish the task of explainability we map relevance from the output to the input features

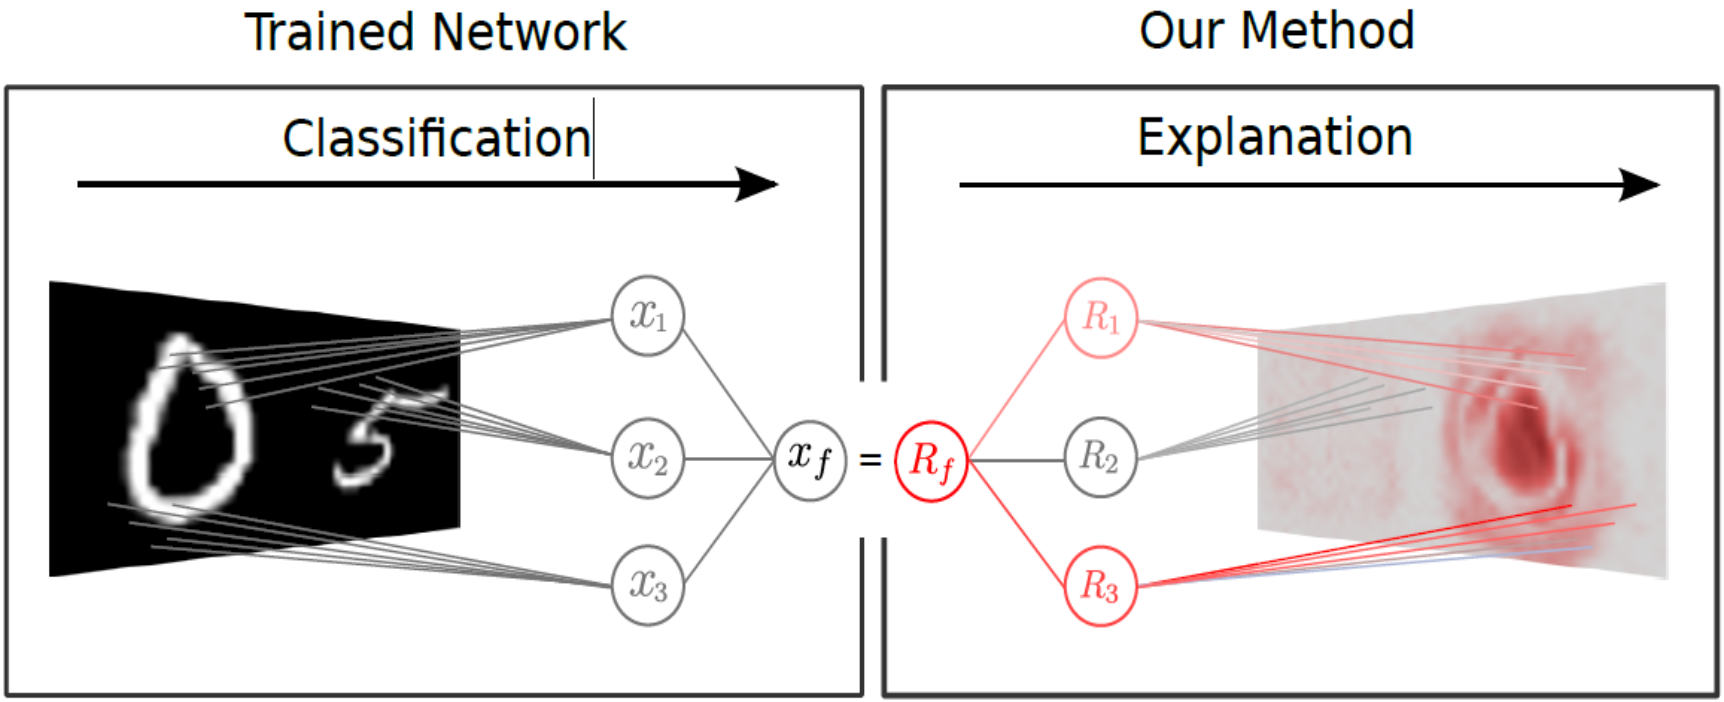

Figure 1: Neural Network Detecting 0 while Distrected by 5

Table of Contents

Mathematical Framework

In the context of image classification we define the following mathematical framework

- Positive valued function $f : \mathbb{R}^d \to \mathbb{R}^+$, where the output $f(x)$ defines either the probability that the object is present or the quantity of the object in question

$\to$ $f(x) > 0$ expresses the presence of the object

- Input $x \in \mathbb{R}^d$, composable in a set of pixel values $x = \{x_p\}$
- Relevance score $R_p(x)$ indicating the relevance of each pixel

$\to$ The relevance score can be displayed in a heatmap denoted by $R(x) = \{R_P(x)\}$

Definitions

Definition 1

A heatmapping $R(x)$ is conservative if the sum of assigned relevances in the pixel space corresponds to the total relevance detected by the model, that is

$$\forall x : f(x) = \sum_p R_p(x) \tag{1}$$

Definition 2

A heatmapping $R(x)$ is positive if all values forming the heatmap are greater or equal to zero, that is:

$$\forall x, p : R_p(x) \geq 0 \tag{2}$$

Definitions

Further we test all algorithms for compliance with definition 1 and 2 hence we introduce the definition of consistency which forth on is a measure of correctness of a technique

Definition 3

A heatmapping $R(x)$ is consistent if it is *conservative* and *positive*. That is, it is consistent if it complies with Definitions 1 and 2.

However consistency is not a measure of quality which can be seen easily on the following example which complies with definition 3

$$\forall p : R_p(x) = \frac{1}{d} \cdot f(x),$$

where d denotes the number of pixels

Taylor Expansion

First order taylor expansion at root point $\tilde{x}$

$$f(x) = f(\tilde{x}) + \left(\frac{\partial f}{\partial x}\Big|_{x=\tilde{x}}\right)^T \cdot (x - \tilde{x}) + \epsilon$$

$$= 0 + \sum_p \underbrace{\frac{\partial f}{\partial x_p}\Big|_{x=\tilde{x}} \cdot (x_p - \tilde{x_p})}_{R_p(x)} + \epsilon$$

The challenge of finding a root point

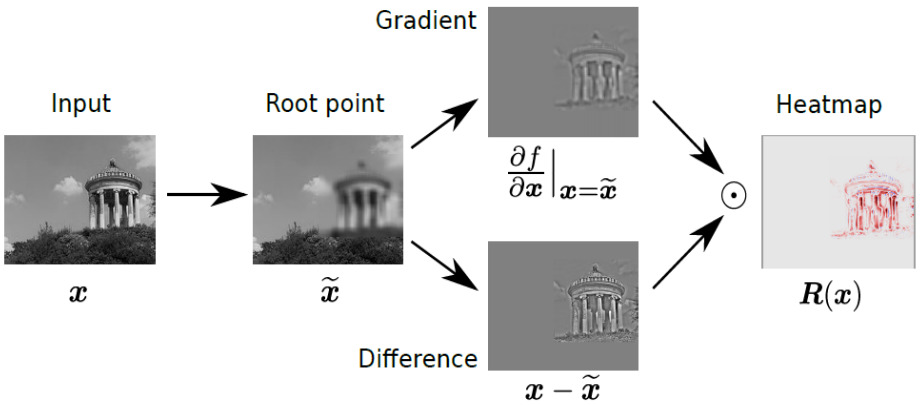

- Potentially more than one root point
- Remove object but deviate as few as possible

$\rightarrow$ $\min_\xi \|\xi - x\|^2$ subject to $f(\xi) = 0$

Sensitivity Analysis

Root point at infinitisimally small distance from the actual point, i.e. $\xi = x - \delta\frac{\partial f}{\partial x}$, where δ is small

If we assume a locally constant function we get

$$
\begin{aligned}
f(x) &= f(\xi) + \left(\frac{\partial f}{\partial x}\Big|_{x=\xi}\right)^T \cdot (x - (x - \delta\frac{\partial f}{\partial x})) + 0 \\
&= f(\xi) + \delta\left(\frac{\partial f}{\partial x}\right)^T \cdot \frac{\partial f}{\partial x} + 0 \\
&= f(\xi) + \sum_p \underbrace{\delta\left(\frac{\partial f}{\partial x}\right)^2}_{R_p} + 0
\end{aligned}
$$

- The heatmap is positive but not conservative
- The heatmap only measures a local effect

Deep Taylor Decomposition

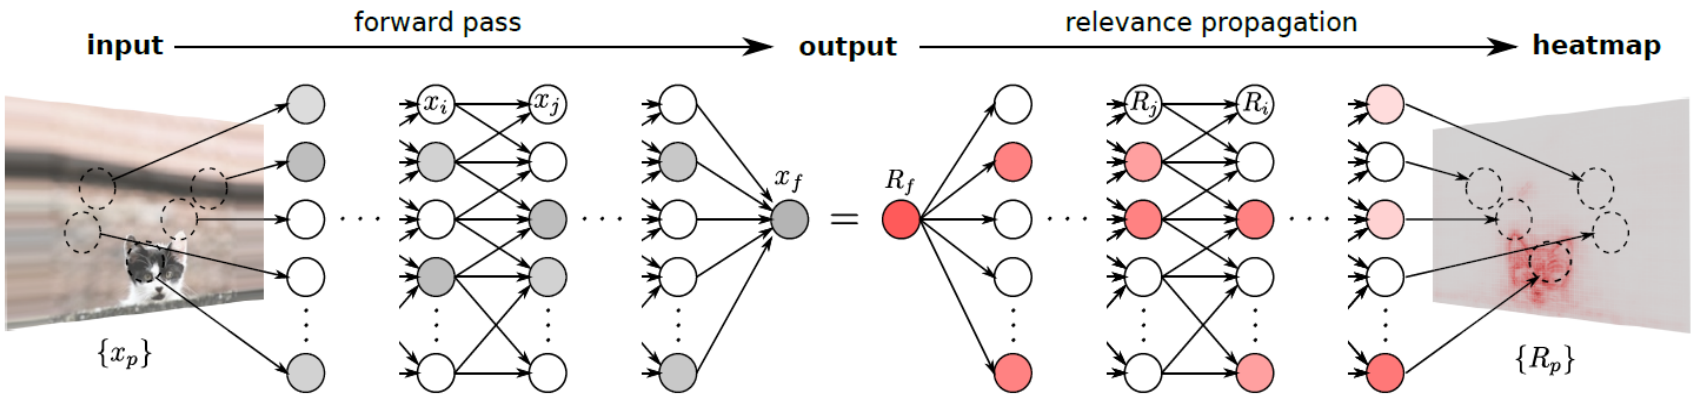

We view each layer as a seperate function and write the Taylor decomposition of $\sum_j R_j$ at $\{x_i\}$ as

$$\sum_j R_j = \left(\frac{\partial(\sum_j R_j)}{\partial\{x_i\}}\Big|_{\{\tilde{x}_i\}} \right)^T \cdot (\{x_i\} - \{\tilde{x}\}) + \epsilon$$

$$= \sum_i \underbrace{\sum_j \frac{\partial R_j}{\partial x_i}\Big|_{\{\tilde{x}\}} \cdot (x_i - \tilde{x}_i)}_{R_i} + \epsilon$$

Deep Taylor Decomposition

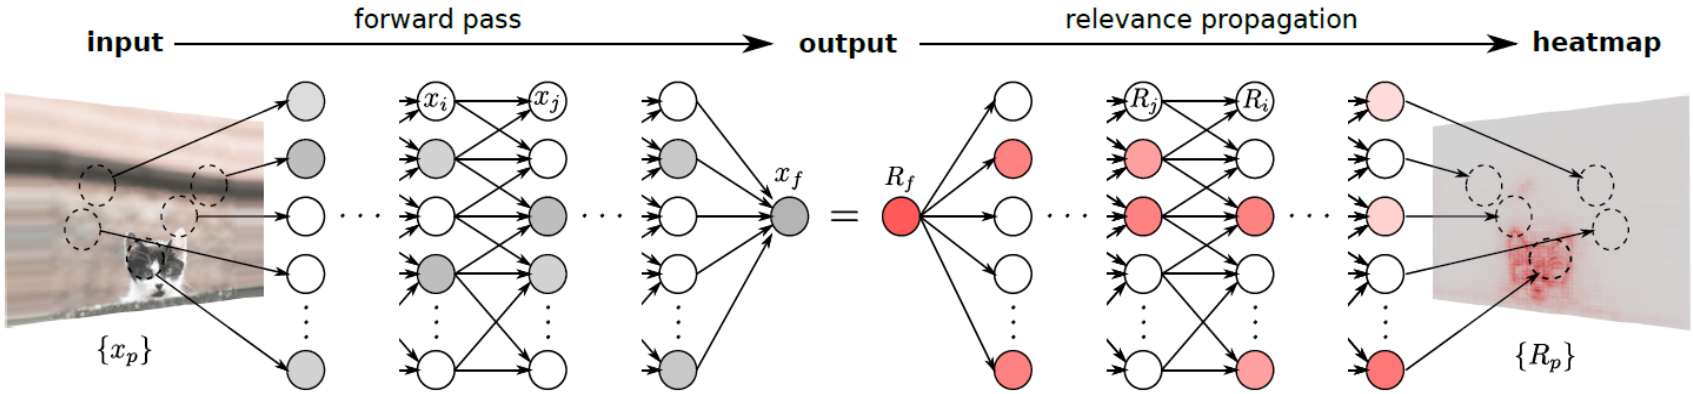

- If each local Taylor decomposition is *conservative* then the chain of equalities is also *conservative* (layer-wise relevance conservation)

$\rightarrow$ $R_f = ... = \sum_i R_i = ... = \sum_p R_p$

- If each local Taylor decomposition is *positive* then the chain of equalities is also *positive*

$\rightarrow$ $R_f, ..., \{R_i\}, ..., \{R_P\} \geq 0$

- If each local Taylor decomposition is *consistent* then the chain of equalities is also *consistent*

Table of Contents

Setting

Consider a simple detection-pooling one layer neural network with

$$x_j = \max(0, \sum_i x_i w_{ij} + b_j)$$

$$x_k = \sum_j x_j, \;\; b_j \leq 0, \forall j$$

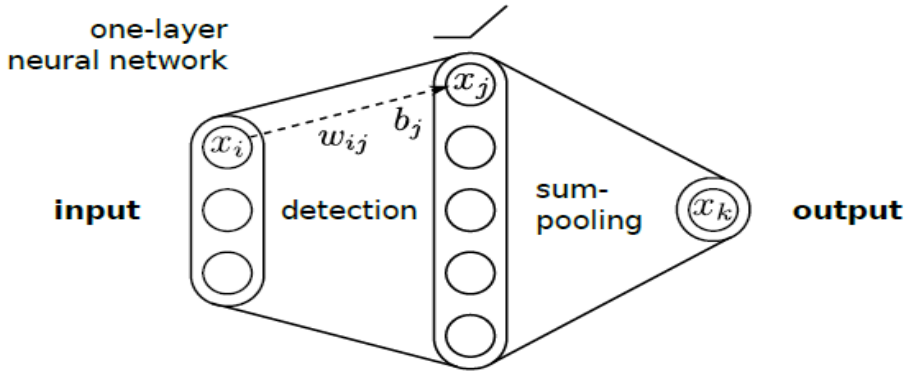

1. $R_k = x_k$ and thus $R_k = \sum_j x_j$
2. Chose a root point and redistribute R_k on neurons x_j
   $\rightarrow R_j = \frac{\partial R_k}{\partial x_j}\Big|_{\{\tilde{x}_j\}} \cdot (x_j - \tilde{x}_j)$, with $\{\tilde{x}_j\} = 0$
3. Since $\frac{\partial R_k}{\partial x_j} = 1$ we obtain $R_j = x_j$
4. Apply Taylor decomposition another time and get
   $R_i = \sum_j \frac{\partial R_j}{\partial x_i}\Big|_{\{\tilde{x}_i\}^{(j)}} \cdot (x_i - \tilde{x}_i^{(j)})$

Derivation of Propagation Rules

Given $R_j = \max(0, \sum_i x_i w_{ij} + b_j)$ and $b_j \leq 0$ and a search direction $\{v_i\}^{(j)}$ in the input space such that

$$\tilde{x}^{(j)} = x_i + t v_i^{(j)} \Leftrightarrow t = \frac{\tilde{x}_i^{(j)} - x_i}{v_i^{(j)}} \tag{3}$$

If the data point itself is not a root point, i.e. $\sum_i x_i w_{ij} + b_j > 0$ the nearest root along $\{v_i\}^{(j)}$ is given by the intersection of equation (3) and $\sum_i \tilde{x}_i^{(j)} w_{ij} + b_j = 0$ which can be resolved to

$$0 = \sum_i x_i w_{ij} + b_j + \sum_i v_i^{(j)} t$$

$$x_i - \tilde{x}_i^{(j)} = \frac{\sum_i x_i w_{ij} + b_j}{\sum_i v_i^{(j)} w_{ij}} v_i^{(j)}$$

Derivation of Propagation Rules

Starting from the Taylor expansion we can plug in

$$x_i - \tilde{x}_i^{(j)} = \frac{\sum_i x_i w_{ij} + b_j}{\sum_i v_i^{(j)} w_{ij}} v_i^{(j)}$$

To get

$$R_i = \sum_j \frac{\partial R_j}{\partial x_i}\Big|_{\{\tilde{x}_i^{(j)}} \cdot (x_i - \tilde{x}_i^{(j)}) = \sum_j w_{ij} \frac{\sum_i x_i w_{ij} + b_j}{\sum_i v_i^{(j)} w_{ij}} v_i^{(j)}$$

$$= \sum_j \frac{v_i^{(j)} w_{ij}}{\sum_i v_i^{(j)} w_{ij}} R_j \tag{4}$$

The relevance propagation rule can be calculated with the following steps

1. Define a segment with search direction $\{v_i\}^{(j)}$
2. The line lies inside the input domain and contains a root point
3. Inject search direction in equation (4)

w^2-rule $\mathcal{X} = \mathbb{R}^d$

Choose a root point which is nearest in the euclidean sense

1. Search direction $\{v_i\}^{(j)} = w_{ij}$ (gradient of R_j)
2. No domain restriction
3. Inject search direction in equation (4)

w^2-rule

$$R_i = \sum_j \frac{\partial R_j}{\partial x_i}\Big|_{\{\tilde{x}_i\}^{(j)}} \cdot (x_i - \tilde{x}_i^{(j)}) = \sum_j \frac{w_{ij}^2}{\sum_{i'} w_{i'j}^2} R_j$$

w^2-rule $\mathcal{X} = \mathbb{R}^d$

Proposition 1

For all functions $g \in G$, the deep Taylor decomposition with the w^2-rule is consistent.

Proof

1. *Conservative*

$$\sum_i R_i = \sum_i \left(\sum_j \frac{w_{ij}^2}{\sum_{i'} w_{i'j}^2} R_j \right)$$

$$= \sum_j \frac{\sum_i w_{ij}^2}{\sum_{i'} w_{i'j}^2} R_j = \sum_j R_j = \sum_j x_j = f(x)$$

2. *Positive*

$$R_i = \sum_j \frac{w_{ij}^2}{\sum_{i'} w_{i'j}^2} R_j = \sum_j \underbrace{w_{ij}^2}_{\geq 0} \underbrace{\frac{1}{\sum_{i'} w_{i'j}^2}}_{>0} \underbrace{R_j}_{\geq 0} \geq 0$$

z^+-rule $\mathcal{X} = \mathbb{R}^d_+$

Search for a root point on the segment $(\{x_i 1_{w_{ij}<0}\}, \{x_i\}) \subset \mathbb{R}^d_+$

1. Search direction $\{v_i\}^{(j)} = x_i - x_i 1_{w_{ij}<0} = x_i 1_{w_{ij}\geq 0}$
2. If $\{x_i\} \in \mathbb{R}^d_+$ so is the whole domain, further for $w_{ij}^- = \min(0, w_{ij})$

$$R_j(\{x_i 1_{w_{ij}<0}\}) = \max(0, \sum_i x_i 1_{w_{ij}<0} w_{ij} + b_j)$$
$$= \max(0, \sum_i x_i w_{ij}^- + b_j) = 0$$

3. Inject search direction in equation (4)

z^+-rule

$$R_i = \sum_j \frac{x_i w_{ij}^+}{\sum_{i'} x_{i'} w_{i'j}^+} R_j$$

z^+-rule $\mathcal{X} = \mathbb{R}^d_+$

Proposition 2

For all functions $g \in G$ and data points $\{x_i\} \in \mathbb{R}^d_+$, the deep Taylor decomposition with the z^+-rule is consistent.

Proof

If $\sum_i x_i w_{ij}^+ > 0$ the same proof as for the w^2-rule applies, if $\sum_i x_i w_{ij}^+ = 0$ it follows that $\forall i : x_i w_{ij} \leq 0$ and

$$R_j = x_j = 0$$

and there is no relevance to redistribute to the lower layers.

z^b-rule $\mathcal{X} = \mathcal{B}$

Often we have a bounded input space $\mathcal{B} = \{\{x_i\} : \forall_{i=1}^d l_i \leq x_i \leq h_i\}$ and a segment $(\{l_i 1_{w_{ij}>0} + h_i 1_{w_{ij}<0}\}, \{x_i\}) \subset \mathcal{B}$

1. Search direction $\{v_i\}^{(j)} = x_i - l_i 1_{w_{ij}>0} - h_i 1_{w_{ij}<0}$
2. If $\{x_i\} \in \mathcal{B}$ so is the whole domain, further for $w_{ij}^- = \min(0, w_{ij})$ and $w_{ij}^+ = \max(0, w_{ij})$

$$R_j(\{l_i 1_{w_{ij}>0} + h_i 1_{w_{ij}<0}\}) = \max(0, \sum_i l_i 1_{w_{ij}>0} w_{ij} + h_i 1_{w_{ij}<0} w_{ij} + b_j)$$
$$= \max(0, \sum_i l_i w_{ij}^+ + h_i w_{ij}^- + b_j) = 0$$

3. Inject search direction in equation (4)

z^b-rule

$$R_i = \sum_j \frac{x_i w_{ij} - l_i w_{ij}^+ - h_i w_{ij}^-}{\sum_{i'} x_{i'} w_{i'j} - l_i w_{i'j}^+ - h_i w_{i'j}^-} R_j$$

z^b-rule $\mathcal{X} = \mathcal{B}$

Proposition 3

For all function $g \in G$ and data points $\{x_i\} \in \mathcal{B}$, the deep Taylor decomposition with the z^b-rule is consistent.

Proof

Since the proof is similar to the proofs of proposition 1 and 2 but lenghty I refer to the literature.

Example MNIST: Setting

Training of a neural network to detect a handwritten digit between 0-3 next to a distractor digit from 4-9 given the following setting:

- images of size 28 x 56 pixels
- 28 x 56 = 1568 input neurons $\{x_i\}$, one hidden layer with 400 neurons $\{x_j\}$ and one output x_k
- weights are random initialized $\{w_{ij}\}$ and bias $\{b_j\}$ is initialized to zero and non negative during training
- Training with 300000 iterations of stochastic gradient descent with a batch size of 20

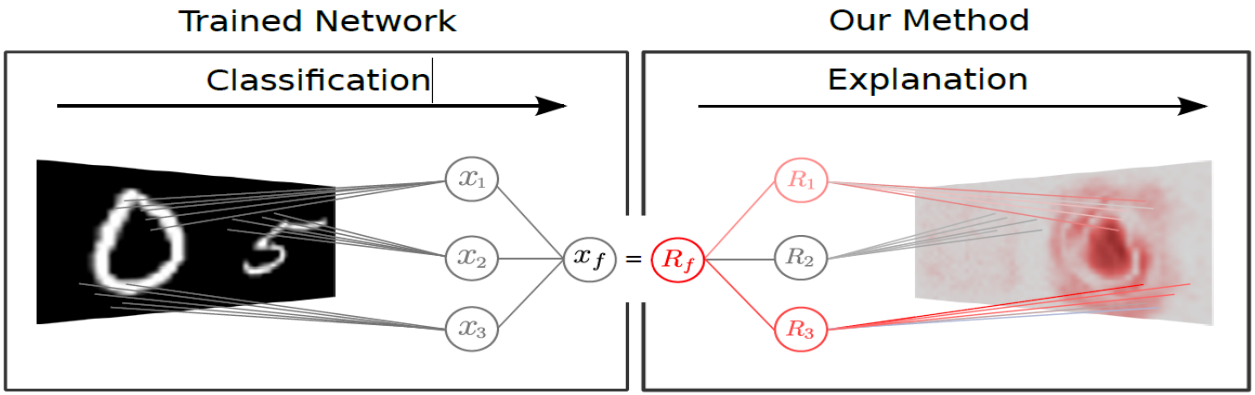

Example MNIST: Heatmaps

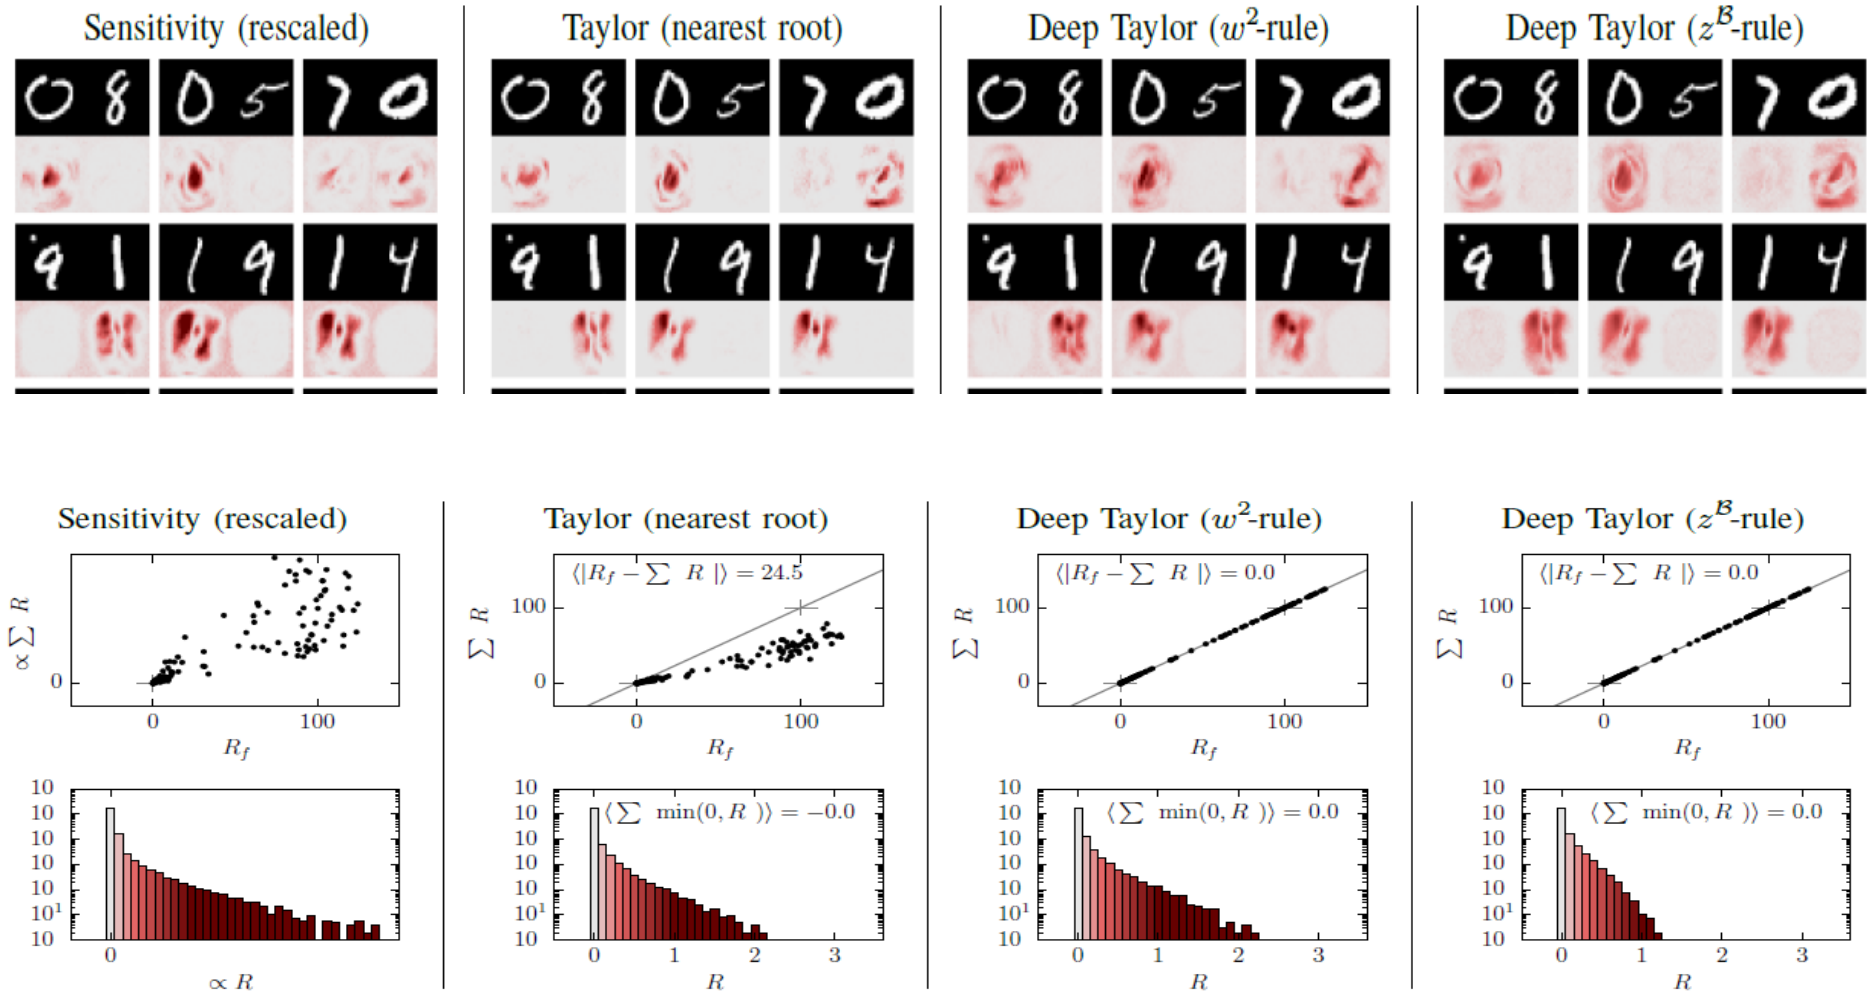

Figure 2: Heatmap and Empirical Results of Consistency

Table of Contents

Deep Networks

Many problems require very complex deep architectures

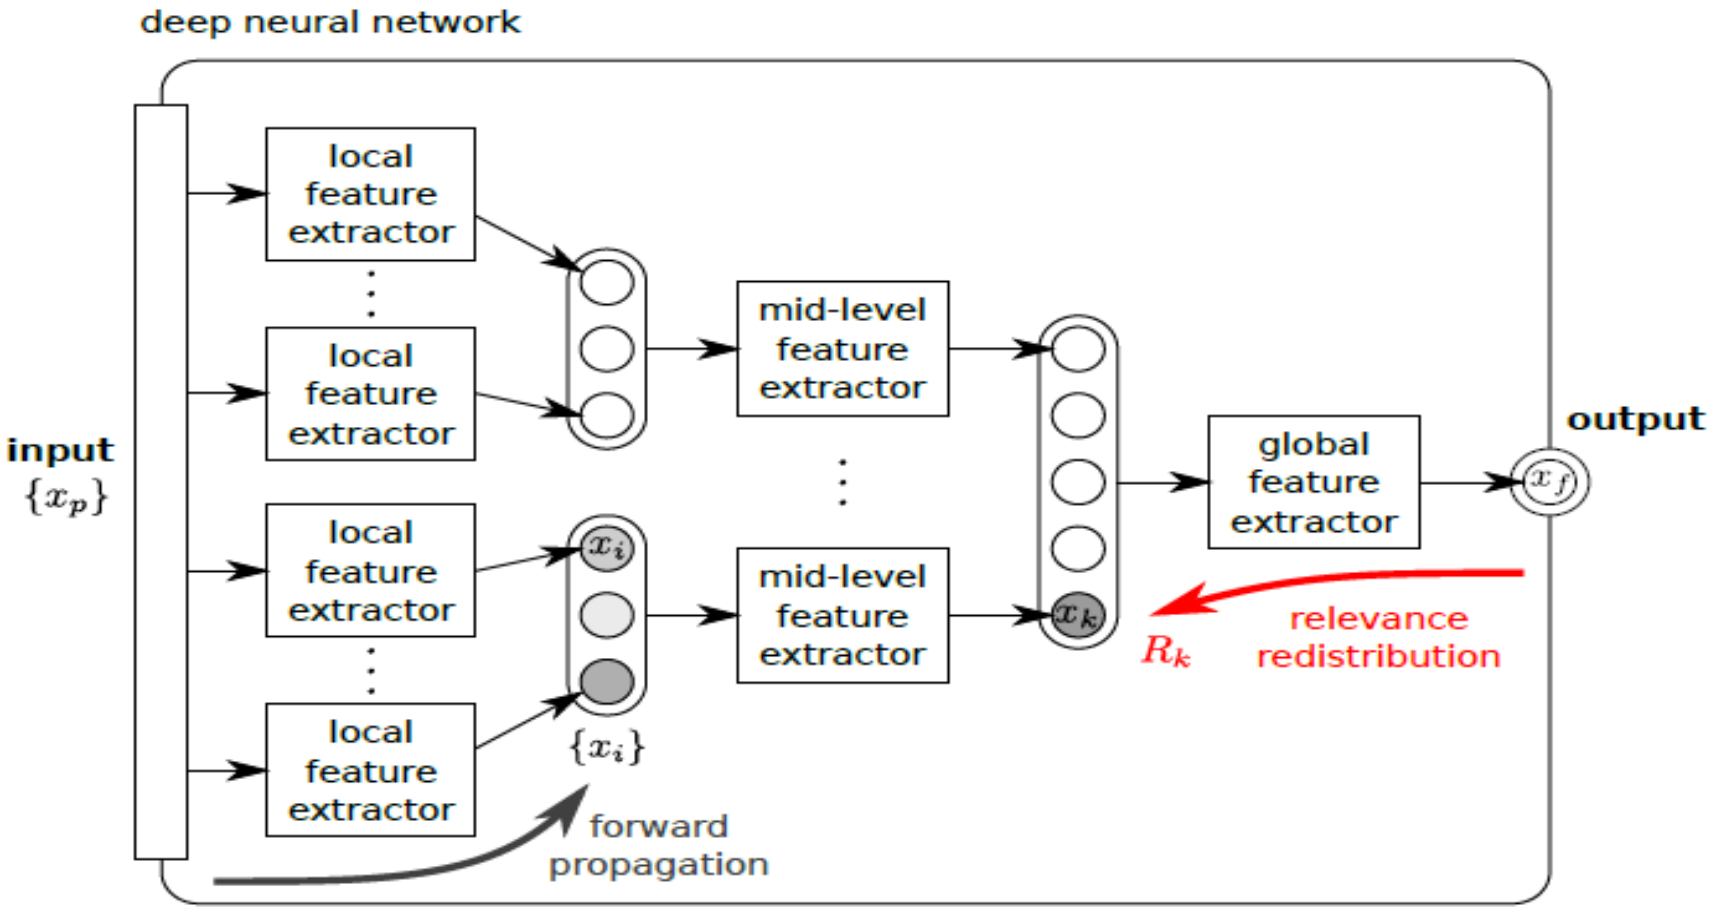

Figure 3: Example Deep Network

Relevance Model

In deep architectures as in figure 3 a mapping between two layers may be unknown, even if it exists

Definition 1

A relevance model is a function that maps a set of neuron activations at a given layer to the relevance of a neuron in a higher layer, and whose output can be redistributed onto its input variables.

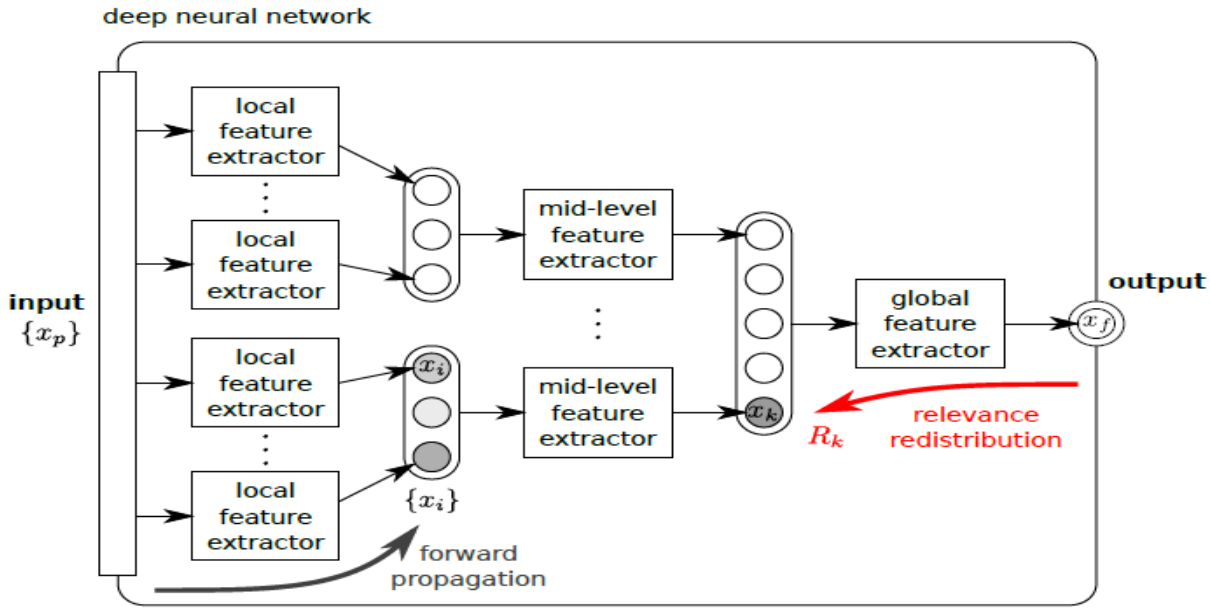

Min-Max Relevance Model

Trainable relevance model designed to incorporate bottom-up and top-down information

$$y_i = \max(0, \sum_i x_i v_{ij} + a_j)$$

$$\hat{R}_k = \sum_j y_j,$$

where $a_j = \min(0, \sum_l R_l v_{lj} + d_j)$ is a negative bias

$\rightarrow$ Compute $\{v_{ij}, v_{lj}, d_j\}$ by minimizing

$$\min \langle (\hat{R}_k - R_k)^2 \rangle$$

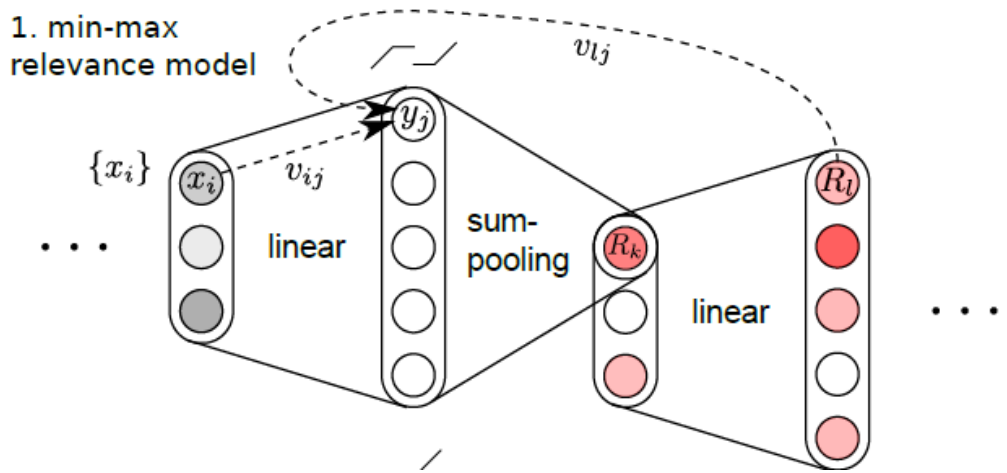

Min-Max Relevance Model

Due to the similar structure we can apply the propagation rules for the one-layer neural network

- Pooling layer

$$R_j = y_j$$

- Detection layer

$$R_i = \sum_j \frac{q_{ij}}{\sum_{i'} q_{i'j}} R_j$$

where $q_{ij} = v_{ij}^2$, $q_{ij} = x_i v_{ij}^+$ or $q_{ij} = x_i v_{ij} - l_i v_{ij}^+ - h_i v_{ij}^-$ for the w^2-rule, z^+-rule and the z^b- rule respectively

$\rightarrow$ The Min-Max relevance model is due to the minimization only approximately consistent

Training-Free Relevance Model

Consider the original network structure

$$x_j = \max(0, \sum_i x_i w_{ij} + b_j)$$

$$x_k = \|\{x_j\}\|_p$$

If the upper layer was explained by the z^+-rule, relevance R_k can be written as

$$R_k = \sum_l \frac{x_k w_{kl}^+}{\sum_{k'} x_{k'} w_{k'l}^+} R_l$$

$$R_k = (\sum_j x_j) \cdot \frac{\|\{x_j\}\|_p}{\|\{x_j\}\|_1} \sum_l \frac{w_{kl}^+ R_l}{\sum_{k'} x_{k'} w_{k'l}^+}$$

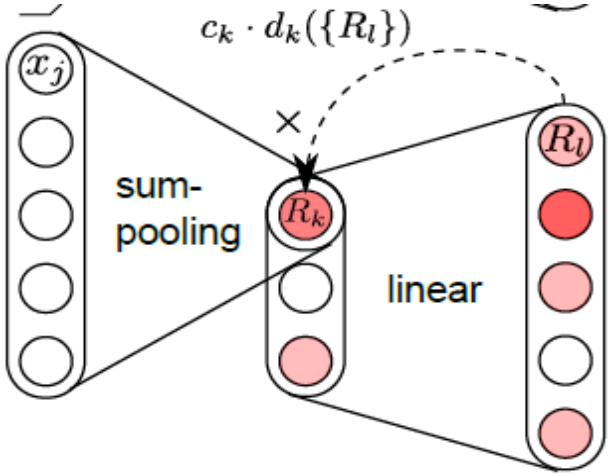

Training-Free Relevance Model

As before we can apply the propagation rules for the one-layer neural network

- Pooling layer

$$R_j = \frac{x_j}{\sum_{j'} x_{j'}} R_k$$

- Detection layer

$$R_i = \sum_j \frac{q_{ij}}{\sum_{i'} q_{i'j}} R_j$$

where $q_{ij} = v_{ij}^2$, $q_{ij} = x_i v_{ij}^+$ or $q_{ij} = x_i v_{ij} - l_i v_{ij}^+ - h_i v_{ij}^-$ for the w^2-rule, z^+-rule and the z^b- rule respectively

$\rightarrow$ The training-free relevance model is consistent

When using the training-free model for the whole network all but the first layer need to be decomposed using the z^+-rule

Table of Contents

Relevance Distribution on the Titanic Dataset

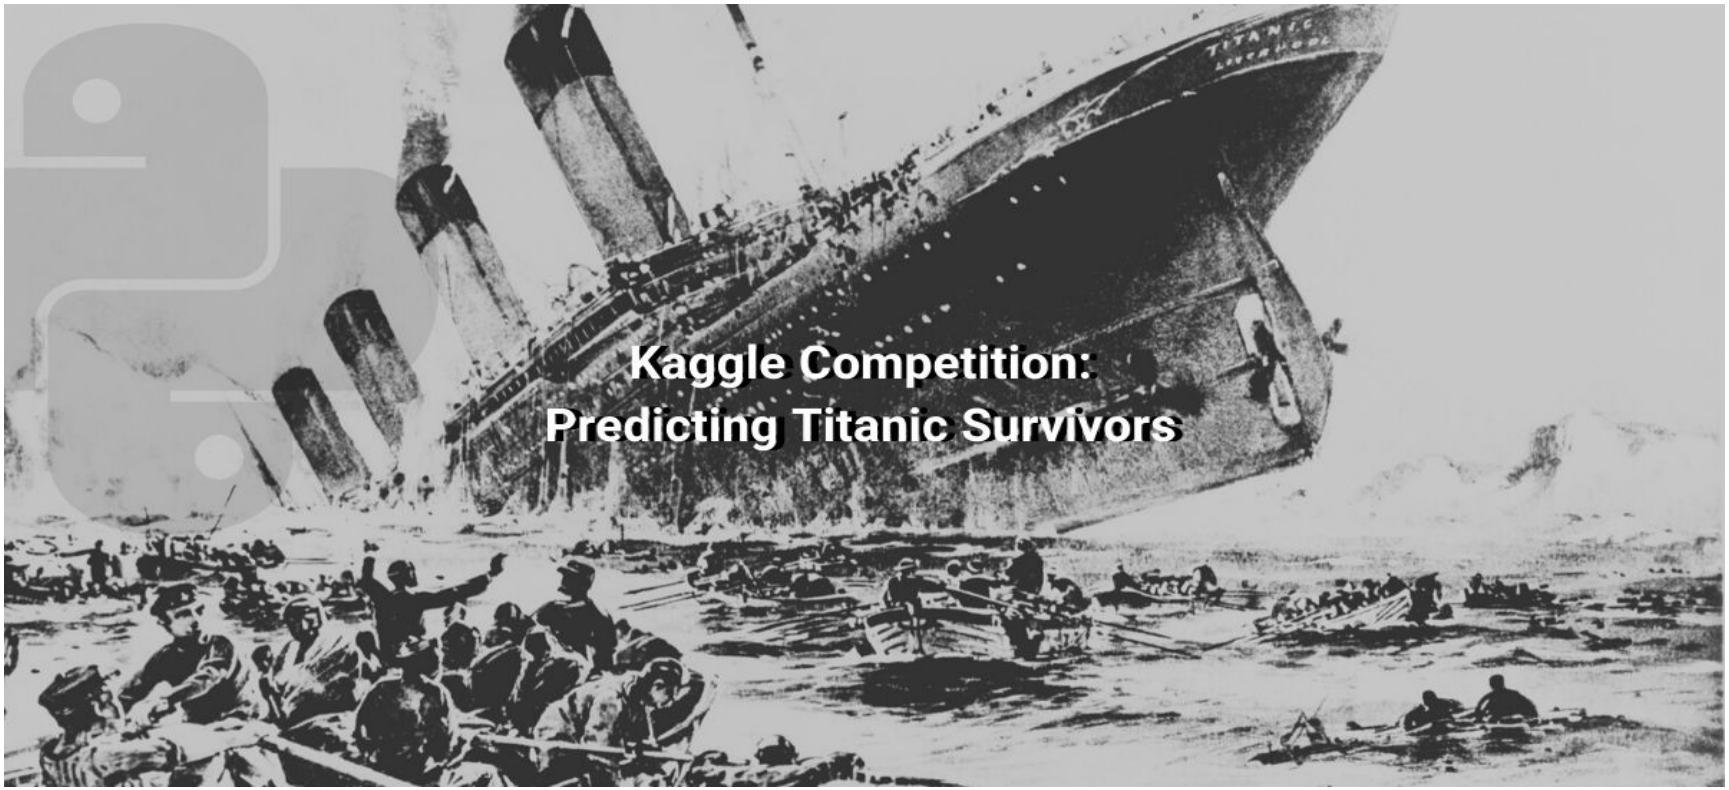

https://github.com/mpommer/Deep-Taylor-Decomposition-Python

References I

John P. Eaton, Charles A. Haas, and John Maxtone-Graham. *Titanic: Triumph and tragedy*. 2nd ed. Nr. Yeovil, Somerset: Patrick Stephens, 1998. ISBN: 185260493X.

Yann LeCun, Corinna Cortes, and CJ Burges. "MNIST handwritten digit database". In: *ATT Labs [Online]. Available: http://yann.lecun.com/exdb/mnist* 2 (2010).

Maximilian Alber et al. "iNNvestigate Neural Networks!" In: *Journal of Machine Learning Research* 20.93 (2019), pp. 1–8. URL: http://jmlr.org/papers/v20/18-540.html.

Grégoire Montavon et al. "Explaining NonLinear Classification Decisions with Deep Taylor Decomposition". In: *Pattern Recognition* 65.2 (2017), pp. 211–222. ISSN: 00313203. DOI: 10.1016/j.patcog.2016.11.008. URL: http://arxiv.org/pdf/1512.02479v1.

References II

Olga Russakovsky et al. "ImageNet Large Scale Visual Recognition Challenge". In: *International Journal of Computer Vision* 115.3 (2015), pp. 211–252. ISSN: 0920-5691. DOI: 10.1007/s11263-015-0816-y.